



ENERGY CONSUMPTION FORECASTING

GROUP 7 : Saniya Khan | Sasidhar Sirivella | Pawan Punera

INTRODUCTION

Machine learning models like Time series produce accurate energy consumption forecasts and they can be used by facilities managers, utility companies and building commissioning projects to implement energy-saving policies. We believe that efforts towards estimating energy consumption and developing tools for researchers to advance their research in energy consumption are necessary for a more scalable and sustainable future.

DATA OVERVIEW

The dataset is obtained from PJM Interconnection which is a regional transmission organization in the United States. PJM is part of the Eastern Interconnection grid operating an electric transmission system serving all or parts of Delaware, Illinois, Indiana, Kentucky, Maryland, Michigan, New Jersey, North Carolina, Ohio, Pennsylvania, Tennessee, Virginia, West Virginia and the District of Columbia. The hourly power consumption data are in megawatt. Here, we are just selecting the power consumption for the Duquesne Light Company, which operates primarily in Pittsburgh and surrounding areas for our project.

DATA DESCRIPTION

Data set contains complete power consumption hourly data through 2005 Dec – 2018 Jan. The file has 119069 observations (hourly data) with two variables as shown below:

Date (type time): time frame at which energy was consumed.

Megawatt Energy consumption (type integer): Energy consumption of a particular region.

PROBLEM STATEMENT


We seek to analyze the electricity consumption data and present the performance of basic forecasting models, STL (Seasonal and Trend decomposition using Loess) with multiple seasonal periods, ETS, ARIMA, TBATS and compare them.

EXPLORATORY ANALYSIS

Our data set has complete **hourly** power consumption data of Duquesne Electric Company from 2005 to 2018. It has 119068 observations of 2 variables (Datetime and DUQ_MW).


However, higher frequency time series often exhibit more complicated seasonal patterns. Hourly data usually has three types of seasonality: a daily pattern, a weekly pattern, and an annual pattern. To deal with such series, we will use the **msts** class which handles multiple seasonality time series which will be discussed later. This allows us to specify all of the frequencies that might be relevant. It is also flexible enough to handle non-integer frequencies.

Before we start exploring the data, we are going to read the data, manipulate it and then visualize our data. Overall, there is a clear trend and strong seasonality in the data set and can be seen in below graph. Now let's visualize for one particular year to better understand the data, its trend and seasonality. Let's consider the power consumption in 2017.



Power consumption peaks during summer months from Mid-June to October and reduces from September and again increased slightly from December. This shows that people use electricity more in summers and winters for cooling and heating.

Next, for the ease of visualization and minimizing the processing time, we are restricting our data to 2013-2017. We then, estimated the trend component and seasonal component of our subset data using `decompose()` function. Trend, seasonal, and irregular components of our data can be estimated using this function.



Decomposition:

- The first row shows our original time series.
- The second panel shows the seasonal component, with the figure being computed by taking the average for each time unit over all periods and then centering it around the mean.
- The third panel plots the trend component and we see a clear trend pattern. This might be sourced from uncaptured extra seasonality from higher natural period in this case and with our huge data. Hence it can be considered as multi-seasonal data. To deal with such series, we will use the `msts` class which handles multiple seasonality time series. This allows us to specify all of the


frequencies that might be relevant. Additionally, using msts instead of ts allows us to specify multiple seasons/cycles, for instance hourly as well as daily: `c(24, 24*7, 24*365.25)`

- The last panel shows the remainder component, which is left over data after removing the trend and seasonal components.

Converting to time series:

The next step is to store the data in a time series object, so that we can use many R functions for analyzing our time series data. To store the data in a time series object, we can use the `ts()` function in R. Sometimes the time series data set that you have, may have been collected at regular intervals that were less than one year, for example, monthly or quarterly. In this case, you can specify the number of times that data was collected per hour by using the frequency parameter in the `ts()` function. Because each row representing a data within hourly interval, we can set `frequency=24`, and we will only use Duquesne Electric Company provider.

Now we see a clearer trend in the below graph after decomposing using `mstl()` and could confirm the daily, weekly and yearly seasonality for our data. Seasonal 168 panel shows the weekly seasonality and seasonal 24 shows the daily seasonality.



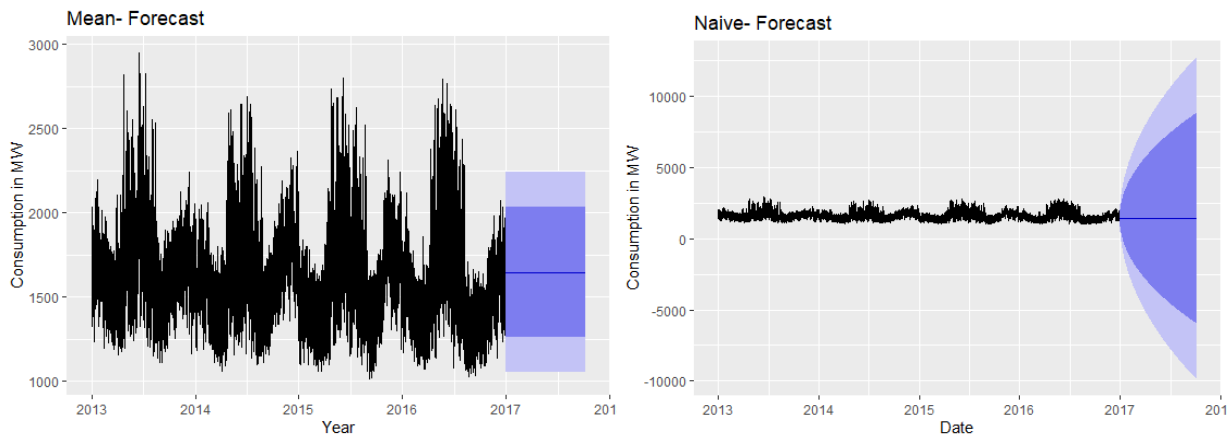
MODELLING – FORECASTING

We first split our data into test and train sets. Train data set includes the data from 2013 -2016 and test data contains data from Jan 2017 -Sept 2017 (around **20 %** of data). To deal with multiple seasonality's, we plan to use ARIMA, TBATS, STLM along with few simple basic forecasting models. Before we jump into ARIMA, STLM and TBATS we will first establish baseline forecasting using simple models like Mean, Naïve, and Seasonal Naïve.

Mean & Naïve Forecasts:

The easiest rough estimate for any forecast would be simply the mean or naïve models. So, after training the models with train data, we have established the mean & naïve baseline forecasts for the test data set.


Naïve model has very high prediction intervals which makes it a worse model for our data. Although, mean forecasts has low intervals it failed to consider the high seasonality in the data.



SNaive & MSTL (STL + ETS) Forecasts:

SNaive method is useful for highly seasonal data. In this case, we set each forecast to be equal to the last observed value from the same season of the year.

The `mstl()` function is a variation on `stl()` designed to deal with multiple seasonality. It will return multiple seasonal components, as well as a trend and remainder component and here's how it works.




ARIMA & TBATS Forecasts:

With multiple seasonality's, we can use Fourier terms and will fit a dynamic harmonic regression model with **ARIMA**. The only drawback here is that it assumes the frequencies stays constant.

A TBATS model differs from dynamic harmonic regression in that the seasonality is allowed to change slowly over time in a TBATS model, while harmonic regression terms force the seasonal patterns to repeat periodically without changing. One drawback of TBATS models, however, is that they can be slow to estimate, especially with long time series. One advantage of the TBATS model is the seasonality is allowed to change slowly over time.

Here the prediction intervals appear to be much too wide – something that seems to happen quite often with TBATS models unfortunately.



CONCLUSION

Judgement Criteria: We are going to use the RMSE, MAE & size of the prediction intervals as the metrics to compare different models.

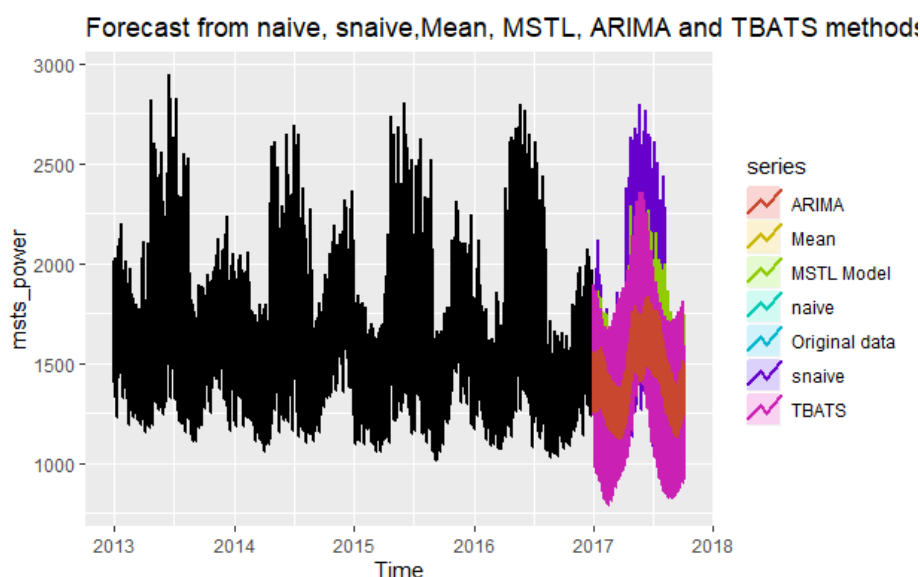
Forecast Metrics Comparison:

```
> kable(summary_table)
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	split	Method
	0.0000000	302.64348	237.40073	-3.2550468	14.6892792	4.6637796	0.9728778	Train	Mean
	-84.6340688	303.59034	245.62702	-8.8493055	16.5967126	4.8253865	NA	Test	Mean
	-0.0040527	70.47619	50.90308	-0.0952334	3.1421760	1.0000000	0.6255802	Train	Naive
	140.8913743	323.81252	238.05178	6.0598402	14.1099969	4.6765697	NA	Test	Naive
	0.0000000	70.47619	50.90313	-0.0949793	3.1421745	1.0000010	0.6255802	Train	Drift
	154.1195266	328.50970	242.70399	6.9472810	14.3206002	4.7679631	NA	Test	Drift
	-32.0274046	354.41441	271.53663	-4.3722712	16.9088447	5.3343855	0.9620306	Train	Snaive
	-53.1647005	436.71906	345.04765	-6.2877663	22.4494717	6.7785225	NA	Test	Snaive
	-0.0046539	25.73499	15.61940	-0.0116438	0.9665696	0.3068460	0.0328774	Train	STLM
	3.3903393	421.15475	336.70676	-3.7400250	21.7411500	6.6146643	NA	Test	STLM
	0.4134288	48.33397	32.73674	-0.0468720	2.0272194	0.6431191	0.0337125	Train	ARIMA
	140.2493658	426.66738	325.07048	5.0128208	19.8811148	6.3860675	NA	Test	ARIMA
	0.0069858	44.05777	27.75356	-0.0416685	1.7313406	0.5452235	0.0129307	Train	TBATS
	103.3999728	464.34704	369.02577	2.9842944	23.4373714	7.2495770	NA	Test	TBATS

Looking at forecast from all the models, the forecast from mstl() or stlm is showing a better performance (**lower RMSE & smaller prediction intervals**) and is definitely our winner here.

Comparing all the forecasts with Plots:



REFERENCES

Data Source. (2018). Retrieved from https://www.kaggle.com/robikscube/hourly-energy-consumption?select=DUQ_hourly.csv

Rob J Hyndman and George Athanasopoulos. (2018). Forecasting: Principles and Practice. Retrieved from <https://otexts.com/fpp2/complexseasonality.html>

Ajeng Prastiwi. (2019, June 23). Multiple Seasonality. Retrieved from <https://rpubs.com/AlgoritmaAcademy/multiseasonality>

Bhide, A. (2019). Reference. Retrieved from <https://www.kaggle.com/apoorvabhide/energy-consumption-time-series-forecasting-in-r/#data>

APPENDIX

Step wise results from R:

Loading the data:

```
## Reading the data
duq_pc <- read.csv(file.choose(),stringsAsFactors = F)
> ## Reading the data
> head(duq_pc)
  Datetime DUQ_MW
1 2005-12-31 01:00:00 1458
2 2005-12-31 02:00:00 1377
3 2005-12-31 03:00:00 1351
4 2005-12-31 04:00:00 1336
5 2005-12-31 05:00:00 1356
6 2005-12-31 06:00:00 1372
> str(duq_pc)
'data.frame': 119068 obs. of 2 variables:
 $ Datetime: POSIXct, format: "2005-12-31 01:00:00" "2005-12-31 02:00:00" ...
 $ DUQ_MW : num 1458 1377 1351 1336 1356 ...
>
```

Subsetting and converting the data into a time series object

```
##-----Modelling-----
#For ease of visualisation and minimising the processing time, we are restricting our
#data to 2011-2017

duq_new <- duq_pc[duq_pc$Datetime >= '2013-01-01 00:00:00' & duq_pc$Datetime <= '2017-09-30 00:00:00',]
#Dividing our data into train and test

duq_train <- duq_new[duq_new$Datetime <= '2016-12-31',]
duq_test <- duq_new[duq_new$Datetime >= '2017-01-01',]
msts_power <- msts(duq_train$DUQ_MW, seasonal.periods = c(24,24*7,24*365.25), start = decimal_date(as.POSIXct("2013-01-01 00:00:00")))
```

Mean Forecasting:

```
> summary(mean_baseline)

Forecast method: Mean

Model Information:
$mu
[1] 1646.525

$mu.se
[1] 1.61682

$s.d
[1] 302.6478

$bootstrap
[1] FALSE

$call
meanf(y = msts_power, h = 24 * 7 * 40)

attr(,"class")
[1] "meanF"

Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 3.508914e-14 302.6435 237.4007 -3.255047 14.68928 0.8742862 0.9728778

Forecasts:
      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
2016.99715    1646.525 1258.654 2034.397 1053.318 2239.733
2016.99726    1646.525 1258.654 2034.397 1053.318 2239.733
2016.99738    1646.525 1258.654 2034.397 1053.318 2239.733
2016.99749    1646.525 1258.654 2034.397 1053.318 2239.733
2016.99760    1646.525 1258.654 2034.397 1053.318 2239.733
2016.99772    1646.525 1258.654 2034.397 1053.318 2239.733
```

STLM Forecasting (STL+ETS):

```
> summary(fcast_mst1)

Forecast method: STL + ETS(A,Ad,N)

Model Information:
ETS(A,Ad,N)

Call:
ets(y = na.interp(x), model = etsmodel, allow.multiplicative.trend = allow.multiplicativ
e.trend)

Smoothing parameters:
  alpha = 0.9999
  beta  = 0.058
  phi   = 0.8

Initial states:
  l = 1754.1249
  b = 6.0828

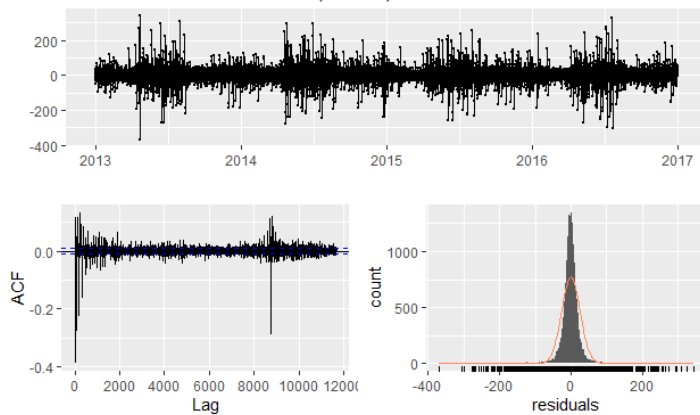
sigma: 25.7368

      AIC      AICC      BIC
594270.6 594270.6 594321.4

Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set -0.004653861 25.73499 15.6194 -0.0116438 0.9665696 0.05752227 0.03287738

Forecasts:
      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
2016.99715    1313.155 1280.173 1346.138 1262.7121 1363.599
2016.99726    1286.233 1238.4963 1333.970 1213.2258 1359.241
2016.99738    1278.740 1219.1098 1338.370 1187.5435 1369.936
2016.99749    1278.405 1209.3795 1348.421 1171.3089 1385.500
```

Residuals from STL + ETS(A,Ad,N)



Naïve Forecasting :

```
> summary(fcast_naive)
```

Forecast method: Naive method

Model Information:
Call: naive(y = msts_power, h = 24 * 7 * 40)

Residual sd: 70.4772

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Training set	-0.004052743	70.47619	50.90308	-0.09523341	3.142176	0.187463	0.6255802

Forecasts:

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2016.99715	1421	1330.6811	1511.319	1282.869205	1559.131
2016.99726	1421	1293.2698	1548.730	1225.653556	1616.346
2016.99738	1421	1264.5631	1577.437	1181.750445	1660.250
2016.99749	1421	1240.3623	1601.638	1144.738410	1697.262
2016.99760	1421	1219.0409	1622.959	1112.130152	1729.870
2016.99772	1421	1199.7648	1642.235	1082.650034	1759.350
2016.99783	1421	1182.0387	1659.961	1055.540268	1786.460
2016.99795	1421	1165.5397	1676.460	1030.307112	1811.693
2016.99806	1421	1150.0434	1691.957	1006.607614	1835.392
2016.99817	1421	1135.3866	1706.613	984.192072	1857.808
2016.99829	1421	1121.4462	1720.554	962.871980	1879.128
2016.99840	1421	1108.1262	1733.874	942.500889	1899.499
2016.99852	1421	1095.3507	1746.649	922.962335	1919.038
2016.99863	1421	1083.0577	1758.942	904.161890	1937.838
2016.99875	1421	1071.1965	1770.803	886.021731	1955.978
2016.99886	1421	1059.7245	1782.275	868.476819	1973.523
2016.99897	1421	1048.6057	1793.394	851.472141	1990.528
2016.99909	1421	1037.8095	1804.191	834.960668	2007.039

Seasonal Naïve Forecasting:

```
> summary(fcast_snaive)
```

Forecast method: Seasonal naive method

Model Information:
Call: snaive(y = msts_power, h = 24 * 7 * 40)

Residual sd: 352.971

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Training set	-32.0274	354.4144	271.5366	-4.372271	16.90884	1	0.9620306

Forecasts:

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2016.99715	1427	972.7997	1881.2	732.3605	2121.639
2016.99726	1412	957.7997	1866.2	717.3605	2106.639
2016.99738	1400	945.7997	1854.2	705.3605	2094.639
2016.99749	1398	943.7997	1852.2	703.3605	2092.639
2016.99760	1416	961.7997	1870.2	721.3605	2110.639
2016.99772	1458	1003.7997	1912.2	763.3605	2152.639
2016.99783	1417	962.7997	1871.2	722.3605	2111.639
2016.99795	1389	934.7997	1843.2	694.3605	2083.639
2016.99806	1373	918.7997	1827.2	678.3605	2067.639
2016.99817	1366	911.7997	1820.2	671.3605	2060.639
2016.99829	1389	934.7997	1843.2	694.3605	2083.639
2016.99840	1445	990.7997	1899.2	750.3605	2139.639
2016.99852	1549	1094.7997	2003.2	854.3605	2243.639
2016.99863	1613	1158.7997	2067.2	918.3605	2307.639
2016.99875	1651	1196.7997	2105.2	956.3605	2345.639
2016.99886	1684	1229.7997	2138.2	989.3605	2378.639
2016.99897	1696	1241.7997	2150.2	1001.3605	2390.639
2016.99909	1716	1261.7997	2170.2	1021.3605	2410.639
2016.99920	1717	1262.7997	2171.2	1022.3605	2411.639

ARIMA with Dynamic Harmonic Regression:

```
> f_fourier <- forecast(fourier_power, xreg=fourier(msts_power, K=
7*40))
> f_fourier
```

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2016.99715		1366.697	1314.7254	1420.724	1288.0180	1450.183
2016.99726		1318.215	1236.1781	1405.697	1194.8375	1454.333
2016.99738		1284.173	1178.6474	1399.146	1126.3425	1464.119
2016.99749		1261.272	1138.3261	1397.496	1078.1709	1475.468
2016.99760		1250.339	1114.4949	1402.741	1048.6638	1490.800
2016.99772		1259.261	1112.8239	1424.967	1042.3296	1521.340
2016.99783		1285.288	1129.6571	1462.360	1055.0512	1565.768
2016.99795		1318.449	1155.2941	1504.645	1077.2645	1613.631
2016.99806		1349.360	1180.7442	1542.055	1100.1889	1654.964
2016.99817		1389.101	1214.9891	1588.165	1131.8360	1704.843
2016.99829		1439.384	1258.9112	1645.730	1172.7234	1766.680
2016.99840		1479.920	1294.3394	1692.109	1205.7138	1816.486
2016.99852		1509.667	1320.1608	1726.376	1229.6709	1853.418
2016.99863		1536.510	1343.2791	1757.537	1251.0294	1887.136
2016.99875		1558.422	1362.0580	1783.096	1268.3325	1914.861
2016.99886		1559.613	1362.8488	1784.786	1268.9455	1916.862
2016.99897		1537.662	1343.5939	1759.761	1250.9814	1890.039
2016.99909		1513.324	1322.3250	1731.910	1231.1774	1860.129

TBATS Forecasting:

```
> tbats_power
TBATS(0.825, {0,0}, -, {<24,5>, <168,5>, <8766,6>})

call: tbats(y = msts_power)

Parameters
  Lambda: 0.824918
  Alpha: 0.9913571
  Gamma-1 values: 0.02230484 0.09414321 -0.06695658
  Gamma-2 values: 0.01192718 0.01054932 0.01004382
```

```
> f_tbats <- forecast(tbats_power, h = 24*7*40)
> f_tbats
```

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2016.99715		1346.7512	1292.2077	1401.684	1263.4953	1430.919
2016.99726		1300.1211	1216.7904	1384.398	1173.0749	1429.381
2016.99738		1293.2136	1188.6486	1399.282	1133.9320	1456.013
2016.99749		1318.6596	1198.0658	1441.218	1135.0649	1506.853
2016.99760		1366.7435	1233.1844	1502.630	1163.4788	1575.458
2016.99772		1436.3286	1290.8032	1584.486	1214.8953	1663.924
2016.99783		1522.6235	1365.3888	1682.757	1283.3988	1768.637
2016.99795		1604.6794	1436.7551	1775.742	1349.2102	1867.499
2016.99806		1659.3296	1482.5890	1839.434	1390.4753	1936.061
2016.99817		1686.4501	1502.5248	1873.961	1406.7059	1974.592
2016.99829		1706.5178	1516.1912	1900.641	1417.0768	2004.851
2016.99840		1729.5298	1533.3264	1929.716	1431.1833	2037.205
2016.99852		1739.7024	1538.6997	1944.862	1434.0937	2055.049
2016.99863		1721.4846	1517.0929	1930.223	1410.7785	2042.374
2016.99875		1692.2459	1485.2275	1903.804	1377.6106	2017.518
2016.99886		1691.8412	1481.8637	1906.492	1372.7406	2021.892
2016.99897		1737.4215	1523.8235	1955.727	1412.7974	2073.076
2016.99909		1800.3795	1583.2563	2022.195	1470.3543	2141.397
2016.99920		1834.7804	1615.0908	2059.183	1500.8381	2179.764
2016.99932		1822.2859	1601.1417	2048.240	1486.1631	2169.677
2016.99943		1778.7202	1556.9068	2005.495	1441.6440	2127.420
2016.99954		1721.9474	1500.0882	1948.938	1384.8809	2071.038
2016.99966		1651.5430	1430.3129	1878.099	1315.5307	2000.037

Next step, we have compared the accuracies

```
## Comparing the accuracies
mean_results <- accuracy(mean_baseline,duq_test$DUQ_MW)
naive_results <- accuracy(fcast_naive,duq_test$DUQ_MW)
rwf_results <- accuracy(fcast_rwf,duq_test$DUQ_MW)
snaive_results <- accuracy(fcast_snaive,duq_test$DUQ_MW)
stlm_model_results<- accuracy(fcast_mstl,duq_test$DUQ_MW)
arima_results<- accuracy(f_fourier, duq_test$DUQ_MW)
tbats_results<- accuracy(f_tbats,duq_test$DUQ_MW)

summary_table= data.table(rbind(mean_results,naive_results,rwf_results,snaive_results,
stlm_model_results,arima_results,tbats_results))
summary_table[,Split:=c("Train","Test","Train","Test","Train","Test","Train","Test")
summary_table[,Method:=c(rep("Mean",2),rep("Naive",2),rep("Drift",2),rep("Snaive",2),
rep("STLM",2),rep("ARIMA",2),rep("TBATS",2))]
kable(summary_table)

> kable(summary_table)
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Split	Method
	0.0000000	302.64348	237.40073	-3.2550468	14.6892792	4.6637796	0.9728778	Train	Mean
	-84.6340688	303.59034	245.62702	-8.8493055	16.5967126	4.8253865	NA	Test	Mean
	-0.0040527	70.47619	50.90308	-0.0952334	3.1421760	1.0000000	0.6255802	Train	Naive
	140.8913743	323.81252	238.05178	6.0598402	14.1099969	4.6765697	NA	Test	Naive
	0.0000000	70.47619	50.90313	-0.0949793	3.1421745	1.0000010	0.6255802	Train	Drift
	154.1195266	328.50970	242.70399	6.9472810	14.3206002	4.7679631	NA	Test	Drift
	-32.0274046	354.41441	271.53663	-4.3722712	16.9088447	5.3343855	0.9620306	Train	Snaive
	-53.1647005	436.71906	345.04765	-6.2877663	22.4494717	6.7785225	NA	Test	Snaive
	-0.0046539	25.73499	15.61940	-0.0116438	0.9665696	0.3068460	0.0328774	Train	STLM
	3.3903393	421.15475	336.70676	-3.7400250	21.7411500	6.6146643	NA	Test	STLM
	0.4134288	48.33397	32.73674	-0.0468720	2.0272194	0.6431191	0.0337125	Train	ARIMA
	140.2493658	426.66738	325.07048	5.0128208	19.8811148	6.3860675	NA	Test	ARIMA
	0.0069858	44.05777	27.75356	-0.0416685	1.7313406	0.5452235	0.0129307	Train	TBATS
	103.3999728	464.34704	369.02577	2.9842944	23.4373714	7.2495770	NA	Test	TBATS

Finally, we plotted forecasts from all the models:

```
## Comparing the models with help of plots
autoplot(msts_power, series = "Original data") +
  geom_line(size = 1) +
  autolayer(fcast_naive, PI = FALSE, size = 1,
    series = "naive") +
  autolayer(fcast_snaive, PI = FALSE, size = 1,
    series = "snaive") +
  autolayer(fcast_mstl, PI = FALSE, size = 1,
    series = "MSTL Model") +
  autolayer(mean_baseline, PI = FALSE, size = 1,
    series = "Mean") +
  autolayer(f_tbats, PI = FALSE, size = 1,
    series = "TBATS") +
  autolayer(f_fourier, PI = FALSE, size = 1,
    series = "ARIMA") +
  ggtitle("Forecast from naive, snaive, Mean, MSTL, ARIMA and TBATS methods")
```

Forecast from naive, snaive, Mean, MSTL, ARIMA and TBATS method:

